

In questo documento spiego alcune operazioni utili che si possono fare per guardare l'andamento dell'esecuzione dell'applicazione

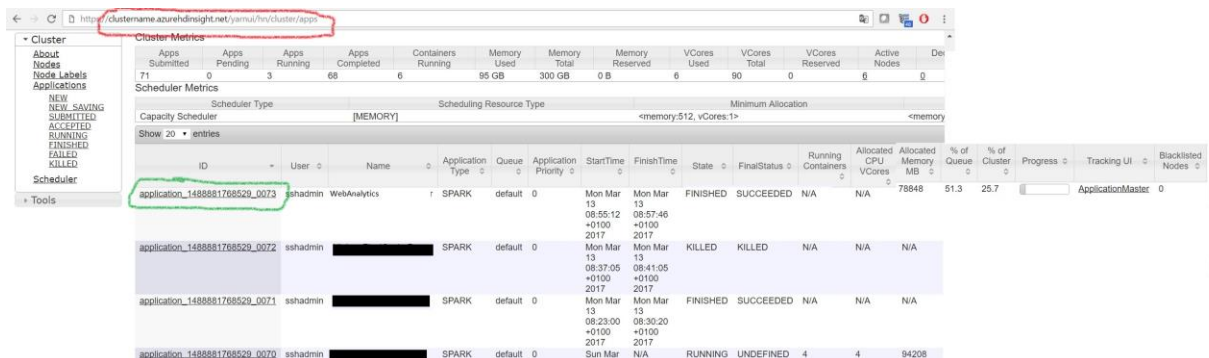
Portarsi sul seguente url:

<https://nomecluster.azurehdinsight.net/yarnui/hn/cluster>

dove nomecluster è il nome del cluster creato da voi

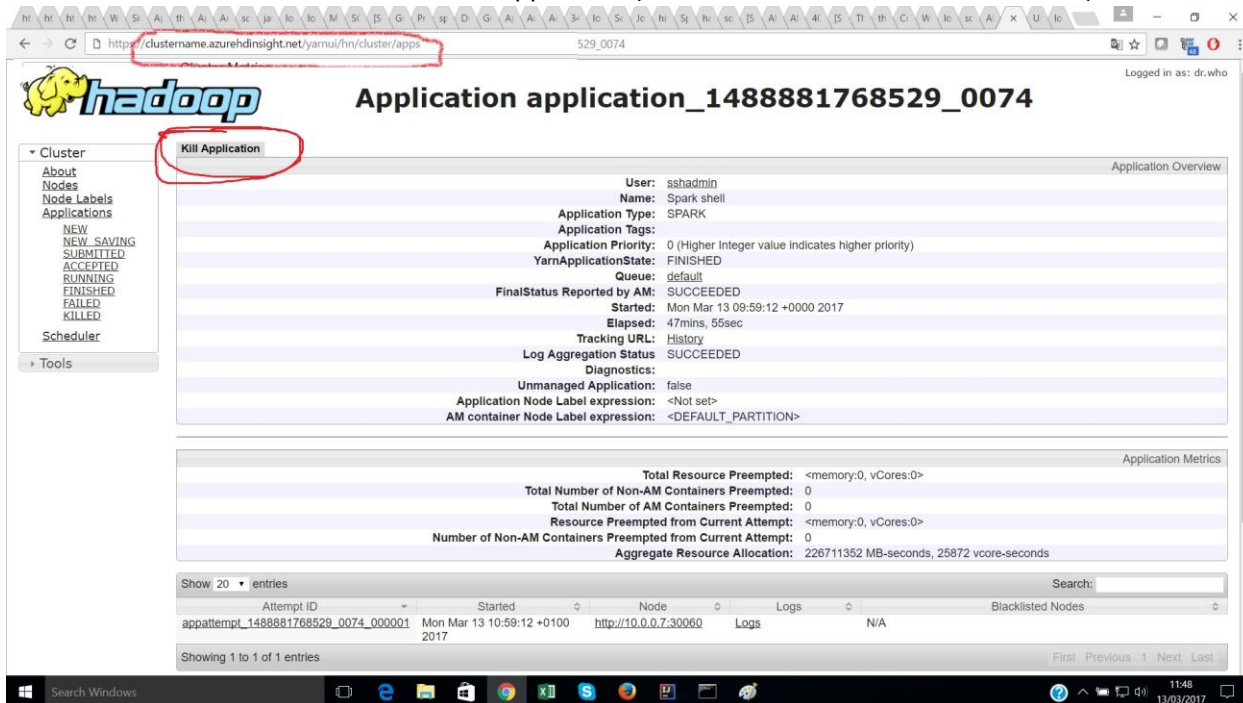
### 1. **KILL APPLICATION** (per voler terminare manualmente l'esecuzione)

Cliccare sul id dell'applicazione che è in esecuzione (sottolineato in verde)



ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1488881768529_0073	sshadmin	WebAnalytics	SPARK	default	0	Mon Mar 13 08:55:12 +0100 2017	Mon Mar 13 08:57:46 +0100 2017	FINISHED	SUCCEEDED	N/A	N/A	78848	51.3	25.7		ApplicationMaster	0
application_1488881768529_0072	sshadmin		SPARK	default	0	Mon Mar 13 08:37:05 +0100 2017	Mon Mar 13 08:41:05 +0100 2017	KILLED	KILLED	N/A	N/A	N/A					
application_1488881768529_0071	sshadmin		SPARK	default	0	Mon Mar 13 08:23:00 +0100 2017	Mon Mar 13 08:30:20 +0100 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A					
application_1488881768529_0070	sshadmin		SPARK	default	0	Sun Mar 12 22:00:00 +0100 2017	N/A	RUNNING	UNDEFINED	4	4	94208					

Successivamente cliccare sulla scritta kill application (anche se non sembra un bottone lo è)



## Application application\_1488881768529\_0074

**Kill Application**

**Application Overview**

User:	sshadmin
Name:	Spark shell
Application Type:	SPARK
Application Tags:	
Application Priority:	0 (Higher Integer value indicates higher priority)
YarnApplicationState:	FINISHED
Queue:	default
FinalStatus Reported by AM:	SUCCEEDED
Started:	Mon Mar 13 09:59:12 +0000 2017
Elapsed:	47mins, 55sec
Tracking URL:	History
Log Aggregation Status:	SUCCEEDED
Diagnostics:	false
Unmanaged Application:	<Not set>
Application Node Label expression:	<Not set>
AM container Node Label expression:	<DEFAULT_PARTITION>

**Application Metrics**

Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	226711352 MB-seconds, 25872 vcore-seconds

Show 20 entries

Attempt ID	Started	Node	Logs	Blacklisted Nodes
appattempt_1488881768529_0074_000001	Mon Mar 13 10:59:12 +0100 2017	http://10.0.0.7:30060	Logs	N/A

Showing 1 to 1 of 1 entries

## 2. Guardare nei Log di Stdout

Per guardare l'esecuzione e i le println che stampo io a video per monitorare l'esecuzione dell'applicazione.

Dalla stessa schermata del punto sopra, cliccare sulla scritta logs (in rosso)

The screenshot shows the Hadoop YARN web interface. The top navigation bar includes 'Cluster', 'Tools', and 'Scheduler'. The main content area is titled 'Application application\_1488881768529\_0070'. It displays various application details such as User (sshadmin), Name (AIAUWrangling), Application Type (SPARK), Application Priority (0), YarnApplicationState (RUNNING), Queue (default), and FinalStatus Reported by AM (Application has not completed yet). Below this, there are Application Metrics showing resource preemption statistics. At the bottom, there is a table of application attempts with columns for Attempt ID, Started, Node, Logs, and Blacklisted Nodes. The 'Logs' column for the first attempt is highlighted with a red circle. The Windows taskbar at the bottom shows the system clock as 12:09 on 13/03/2017.

E cliccare su **stdout** e si visualizzeranno i messaggi visualizzati dalle println

The screenshot shows the Hadoop YARN web interface displaying the stdout logs for container\_e06\_1489404693736\_0014\_01\_000001. The top navigation bar includes 'ResourceManager', 'RM Home', 'NodeManager', and 'Tools'. The main content area is titled 'Logs for container\_e06\_1489404693736\_0014\_01\_000001'. It displays the stdout logs, which include directory information and file lengths. The Windows taskbar at the bottom shows the system clock as 17:04 on 13/03/2017.

3. **Guardare log di errore** di Apache Spark

Nel caso ci fossero degli errori e per vederne i log, dalla schermata precedente cliccare su ***Stderr***, e dal browser con ctrl+F scrivere error e trovare i log di errori.